Section 1 — Project Declaration: Onirium AI Project Charter (V3)

Reference Methodology: PMBOK® Guide Standard **Organization:** Antigravity AI Labs

1. General Project Information

Field	Detail
Project Title	*Onirium AI: Multi-Agent Platform for Symbolic Dream Exploration (V3)*
Primary Sponsor	*Executive Directorate of Antigravity*
Project Lead / AI Architect	*Senior Intelligent Agents Developer – Antigravity*
Last Updated	*June 22, 2026*
Document Status	*Approved with Integrated Visual Architecture Reference*

Section 2 — Purpose and Project Justification

The Onirium AI project arises in response to the growing global demand for well-being, self-knowledge, and entertainment applications powered by Generative Artificial Intelligence. This demand is particularly notable in Latin America and Brazil, where many individuals frequently consult websites or AI models seeking predictive or premonitory interpretations of their dreams. Such practices pose significant risks, including inappropriate access to clinical-type information without expert supervision and exposure to potentially harmful hallucinations generated by large language models.

Historically, dream interpretation has fluctuated between academic clinical psychology (Jungian, Freudian) and localized esoteric cultural folklore. Onirium AI addresses and avoids these risks by automating the narrative and symbolic exploration of dream reports through an ethically grounded, orchestrated multi-agent approach within Antigravity's local ecosystem.

Its fundamental purpose is to provide a playful yet introspective environment that remains safe for users, while structurally mitigating civil-liability risks associated with clinical diagnosis. This is achieved through a rigorous architectural design that embeds ethical and legal guardrails across all system layers.

Section 3 — Project Objectives and Success Criteria

Dimension	Metric / KPI	Success Criterion
Ethical and Legal Safety	Incidents of false negatives in psychological-crisis detection.	0% leakage (absolute blocking of self-harm or psychosis-related language).
AI Performance	Contextual coherence and clinical-hallucination rate.	Below 1% under medical-grade criteria, verified through prompt-audit procedures.
Product and User Retention	MVP user retention and conversational-session frequency.	Positive engagement growth with a limit of 3 free queries per day.

Section 4 — Detailed Project Scope

4.1. Global Scope

The system will initially operate in Spanish, providing full support for the diverse dialectal variants, idioms, and linguistic nuances across Spanish-speaking regions in Latin America, including Mexico, the Andean region, the Caribbean, and the Southern Cone.

4.2. Product Scope

The solution is designed under the development strategy: *"Build a Web Application in Antigravity and Deploy to Cloud Run."* The agent will be accessible exclusively through a responsive, optimized web application and through native or hybrid mobile platforms that consume the same cloud-based infrastructure.

Third-party messaging channels such as WhatsApp or Telegram are explicitly excluded from the MVP, as are voice-cloning components, hyper-realistic avatars, or image-analysis capabilities.

4.3. Context Scope

The agent will maintain long-term persistent memory through an anonymized vector database. This will enable the system to securely correlate dreams narrated at different points in time to identify patterns or "recurring dreams," without storing any personally identifiable information (PII).

Section 5 — Technical Architecture and Multi-Agent Topology in Antigravity

In accordance with best practices from the Antigravity framework and event-driven architecture principles, system processing is distributed across an orchestrated multi-agent topology formally represented in the system's technical blueprint.

5.1. Architectural Model Integration (Official Blueprint)

Infrastructure Technical Note: Due to restrictions in the automated document-provisioning API, binary images must be manually consolidated by the repository administrator using the following verbatim identifier as the official visual asset attached to the Project Charter:

Código

```
======================================================================
[MANDATORY VISUAL REFERENCE BLOCK - AI FLOW ARCHITECTURE]

Required Reference File:
Gemini_Generated_Image_3djf2f3djf2f3djf.png

Description:
Official topological map of Onirium's asynchronous cloud flow.
======================================================================
```

The technical blueprint **Gemini_Generated_Image_3djf2f3djf2f3djf.png** defines the following fixed operational sequence:

- **User → API Gateway & Ethical Safety Layer** — The dream-text input is submitted through the web/mobile application and reaches the security boundary as the first processing step.
- **Pub/Sub Topic (Event Router)** — Routes the request event asynchronously toward the compute backend.
- **Multi-Agent Orchestrator (Cloud Run)** — The central node, deployed as an Antigravity container on Cloud Run, responsible for core logic execution and distribution of tasks to sub-agents.

- **MCP Server Connectors (Model Context Protocol)** — Dynamic context servers (Cultural DB, Symbols, Context) that exchange information bidirectionally with the orchestrator.
- **Vector DB (Pinecone / Weaviate)** — Storage and retrieval of semantic embeddings used to detect recurrence patterns through structured queries.
- **Anonymized BigQuery Dataset (Storage)** — Final, secure analytical repository that receives anonymized outputs for long-term storage and evaluation.

5.2. MCP Server Specifications

- **mcp-server-cultural-folklore** — Provides tools to query regional traditions, myths, and popular Hispanic-American symbolism based on the user's geographic profile.
- **mcp-server-dream-symbology** — Offers indexed access to public-domain dictionaries of psychological archetypes (Jung, Freud), focused exclusively on narrative and personal-growth perspectives, never clinical interpretation.
- **mcp-server-safety-moderation** — Dynamically updates content-filtering rules and provides a geolocated list of humanitarian or psychological support centers.

Section 6 — Ethical Restrictions and International Legal Compliance

6.1. Legal Restrictions

- **Age Control (+18)** — Strict access restriction enforced through a validation gateway during onboarding. Full compliance with COPPA requirements and GDPR child-protection regulations.
- **Data Privacy (GDPR / CCPA)** — Integration of the *Privacy by Design* principle. Mandatory implementation of a **"Complete Vector-History Deletion"** button (Right to Erasure), allowing the user to trigger full deletion with a single click.
- **AI Regulation (EU AI Act)** — Strict adherence to transparency guidelines. The system will proactively inform users that they are interacting with an artificial agent and that all outputs are narrative simulations, not factual truths.

6.2. Safeguards Against Clinical Diagnosis

The system is architecturally prohibited from issuing diagnostic verdicts. Specialist-agent system prompts include explicit lexical blocking. The use of terms such as *"disorder," "depression," "pathology," "schizophrenia,"* or *"trauma"* is strictly forbidden.

If a user introduces mental-health symptomatology, the agent will redirect them to professional resources using a strictly standardized format:

Fixed Safety Disclaimer

"This system is intended solely for entertainment and symbolic exploration. It does not provide clinical evaluation, diagnosis, or therapeutic guidance."

The system is architecturally prohibited from issuing diagnostic verdicts. Specialist-agent prompts include explicit lexical blocking. The use of terms such as *"disorder," "depression," "pathology," "schizophrenia,"* or *"trauma"* is strictly forbidden.

If the user introduces mental-health symptomatology, the agent will redirect them to professional resources using a strictly standardized format:

Fixed Safety Disclaimer

"This system is intended solely for entertainment and symbolic exploration. It does not constitute a substitute for psychological therapy nor does it provide medical diagnoses. If you are experiencing severe emotional distress, please contact a health professional or the following free support lines: [Dynamic_MCP_List]."

Section 7 — User Interface (UI/UX) Design Guidelines

The user interface must evoke an immersive, safe, and mystical environment—visually aligned with the minimalism of modern software architectures while remaining thematically consistent with the subconscious.

- **Tertiary Color Palette (Native Dark Mode)** — The application's primary background will use a deep night-blue tone (#0B0F19). Interactive elements, containers, and active buttons will be highlighted with subtle gradients of nebula violet (#6D28D9) and electric cyan (#06B6D4) to denote active data flows.
- **Anxiety-Reducing Conversational Micro-Interactions** — Agent-generated responses will not appear suddenly as full blocks of text. A progressive word- or phrase-level rendering effect (soft *fade-in*) will be implemented to reduce perceived technical latency and simulate a calm, reflective conversational rhythm.
- **Placement of Legal Disclaimers and Consent Notices** — Age-verification confirmation and acceptance of the data-anonymization policy will be presented in full-screen onboarding screens before allowing the first text input. Additionally, the chat interface will maintain a subtle but readable fixed bottom banner displaying the clinical-responsibility disclaimer.
- **Interactive Visualization of the "Semantic Oniro-Map" Artifact** — Upon completing the interpretation, the interface will offer an alternate tab that transforms the JSON extracted by the Orchestrator Agent into an interactive visual graph. Key symbols will appear as luminous floating nodes interconnected with detected emotions, allowing the user to expand each node and download their **"Dream Logbook"** in a clean Markdown format or as a corporate-style PDF export.

Section 8 — Signatures and Authorizations

By proceeding with the configuration of agent setup files within the local Antigravity environment, the development team agrees to adhere strictly to the ethical and technical boundaries established in this Project Charter.

Section 9 — Technical Annexes

9.1. Orchestrator Prompt (Multi-Agent Governance Example)

This example illustrates how the Multi-Agent Orchestrator governs sub-agents, enforces ethical constraints, and coordinates the full symbolic-interpretation workflow.

Orchestrator Prompt

```
You are the Onirium Multi-Agent Orchestrator.

Your responsibility is to coordinate specialized sub-agents under strict ethical,

legal, and narrative constraints.


SYSTEM RULES:

- Never produce clinical interpretations, diagnoses, or medical terminology.

- Enforce lexical blocking: disorder, depression, pathology, schizophrenia, trauma.

- All outputs must remain symbolic, narrative, and introspective.

- If mental-health symptoms appear, trigger the Safety-Moderation MCP server and

  return the Fixed Safety Disclaimer.


WORKFLOW:

1. Send the user's dream text to:

   - mcp-server-cultural-folklore → cultural symbols

   - mcp-server-dream-symbology → archetypal references

   - mcp-server-context → contextual embeddings

2. Merge all responses into a unified symbolic-narrative interpretation.

3. Generate a structured JSON for the Semantic Oniro-Map.

4. Return a safe narrative interpretation + the JSON artifact.
```

OUTPUT FORMAT:

```
{
  "interpretation": "...",
  "oniromap_json": { ... },
  "safety": null | "triggered"
}
```

9.2. End-to-End Request Lifecycle (Dream Processing Flow)

This section describes the complete lifecycle of a user request within the Onirium AI system.

1. **User Input** The user submits the dream text through the web or mobile interface.
2. **Ethical Safety Layer** The system filters prohibited clinical terminology, self-harm language, or unsafe content.
3. **Pub/Sub Event Router** The request is asynchronously routed to the compute backend.
4. **Multi-Agent Orchestrator**
   - Receives the sanitized input
   - Queries MCP servers
   - Applies ethical guardrails
   - Generates the symbolic interpretation
   - Produces the Oniro-Map JSON artifact
5. **Sub-Agents**
   - Cultural Folklore MCP
   - Dream Symbology MCP
   - Context MCP
   - Safety-Moderation MCP
6. **Vector Database** Embeddings are queried to detect recurring dream patterns.
7. **Output Renderer** The UI displays:
   - symbolic interpretation
   - interactive Oniro-Map
   - legal disclaimers
   - micro-interaction animations

9.3. Technical Glossary

Term	Definition
MCP (Model Context Protocol)	Protocol enabling agents to access external tools such as cultural databases, symbolic dictionaries, and safety modules.
Embedding	Numerical semantic representation of text used to detect similarity and patterns.
Vector Database	Specialized database for storing and querying embeddings (e.g., Pinecone, Weaviate).
Recurring Dream	A pattern detected when multiple dreams share symbols, emotions, or narrative structures.
Oniro-Map	Interactive graph representing symbols, emotions, and relationships extracted from the dream.
Safety Layer	System component that blocks clinical, harmful, or unsafe language.
Orchestrator	Central agent coordinating sub-agents, enforcing rules, and generating final outputs.

9.4. Example JSON — Semantic Oniro-Map (Simplified)

This JSON structure is used by the UI to generate the interactive symbolic graph.

Oniro-Map JSON Example


```
{
  "dream_id": "session_2026_06_22_001",
  "symbols": [
    { "id": "moon", "label": "Moon", "archetype": "intuition", "intensity": 0.82 },
    { "id": "water", "label": "Water", "archetype": "emotion", "intensity": 0.67 },
    { "id": "bridge", "label": "Bridge", "archetype": "transition", "intensity": 0.54 }
  ],
  "emotions": [
    { "id": "calm", "value": 0.71 },
    { "id": "uncertainty", "value": 0.43 }
  ],
```

"connections": [

{ "source": "moon", "target": "calm", "weight": 0.62 },

{ "source": "water", "target": "uncertainty", "weight": 0.48 },

{ "source": "bridge", "target": "uncertainty", "weight": 0.51 }

],

"summary": "The dream suggests an emotional transition guided by intuition and calmness."

}